

# Try Harder 2 Hardening the COREs



Shawn C

# #whois

- \* Shawn C[a.k.a "citypw"]
- \* Day job at TYA infotech
  - \* Open source security consulting
- \* GNU/Linux security engineer
- \* Free/libre SW/FW/HW enthusiasts
- \* Member of EFF/FSF/FSFE/RISC-V
- \* Patient Zer0 at HardenedLinux community(<https://hardenedlinux.github.io/>)



# #whois2

- \* Persmule
- \* Day job at TYA infotech
  - \* Open source security consulting
- \* GNU/Linux engineer
- \* Free/libre SW/FW/HW enthusiasts
- \* Firmware maintainer at HardenedLinux
- \* Old crew at BLUG( <https://beijinglug.club/>)



# #cat /proc/agenda

- \* History: Ring 3
- \* Attacking the core
- \* Under the water
- \* Invincible devil
- \* Hope or delusion?

# #Ring 3

- \* MAC/DAC/Sandboxing( seccomp)
- \* Baseline checks( STIG-4-Debian)
- \* Firewall/IDS/SOC
- \* etc

# #We had a history with Ring 3

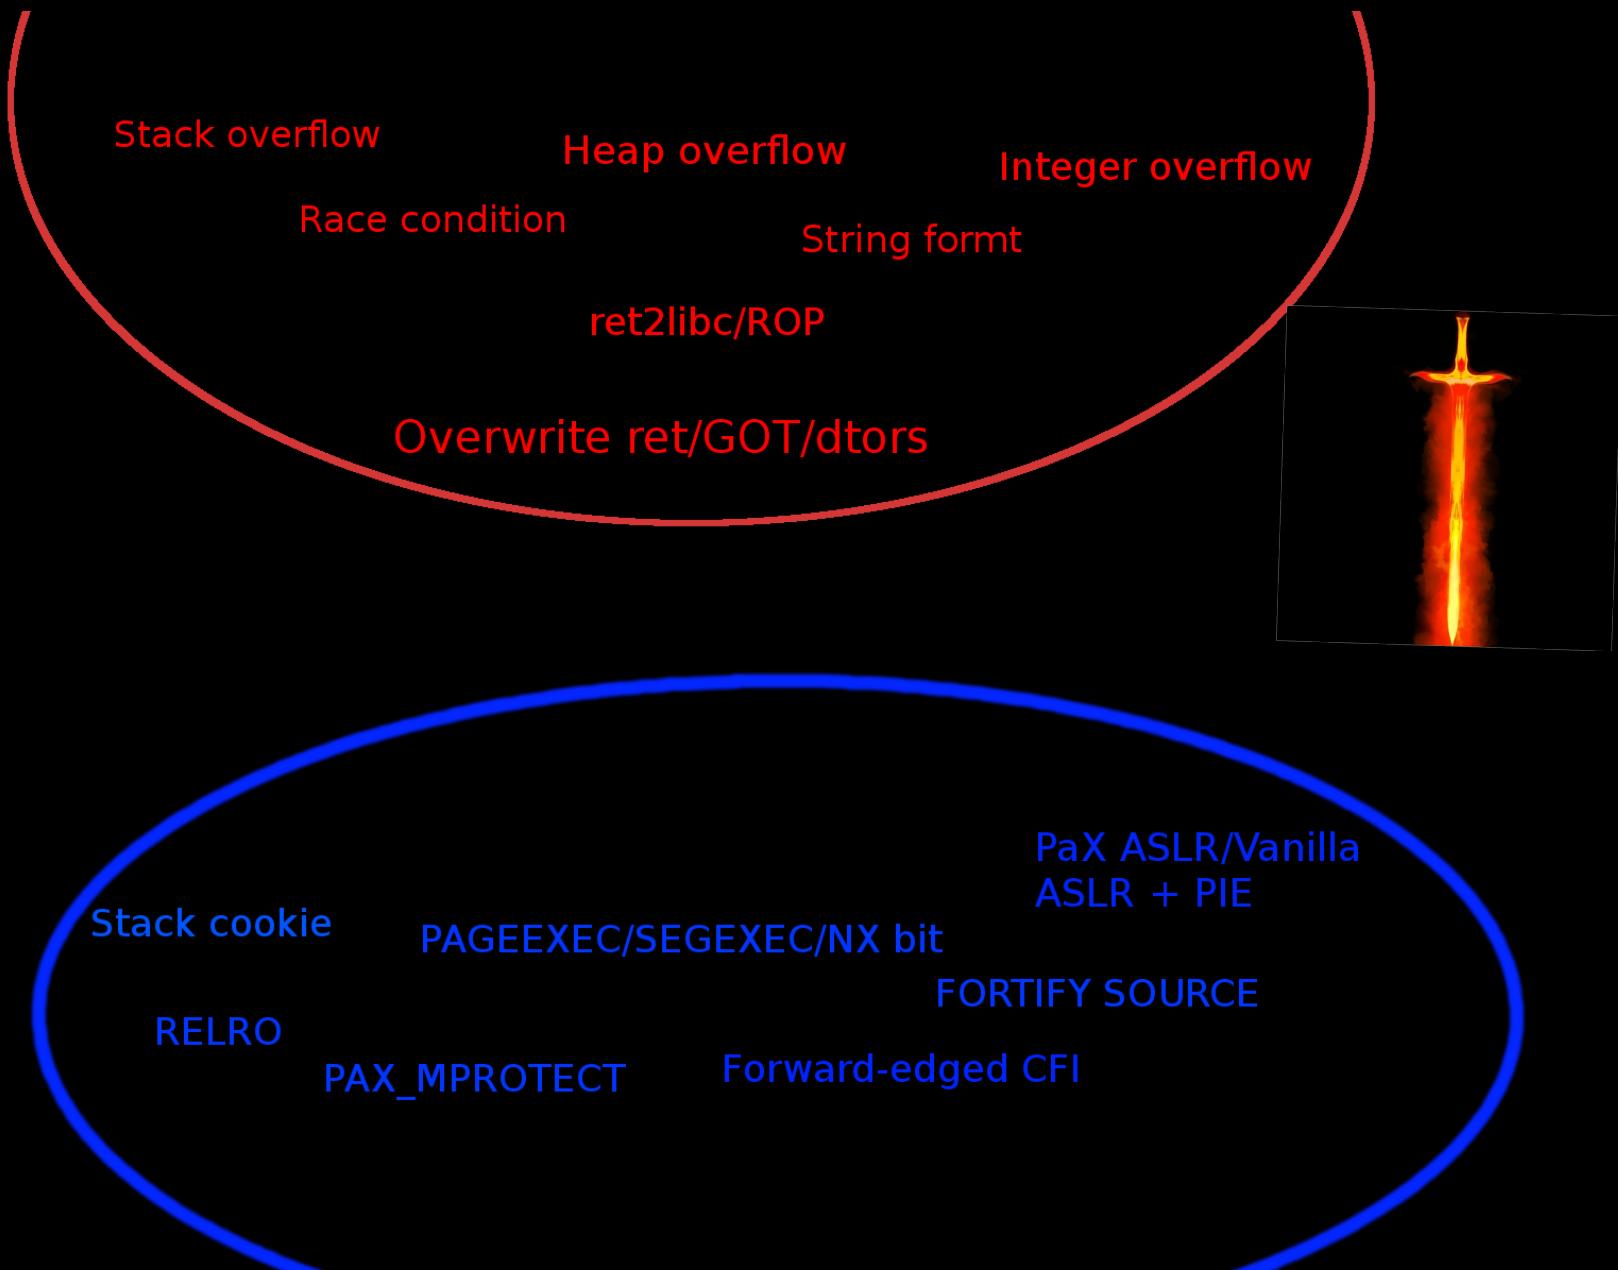
\* Once upon a time, the stack was so "pure"...



- \* Stack/Heap layout
- \* Stack grows down (x86, MIPS)
- \* ESP points to the current top of the stack
- \* EBP points to the current function frame



# #War at Ring 3



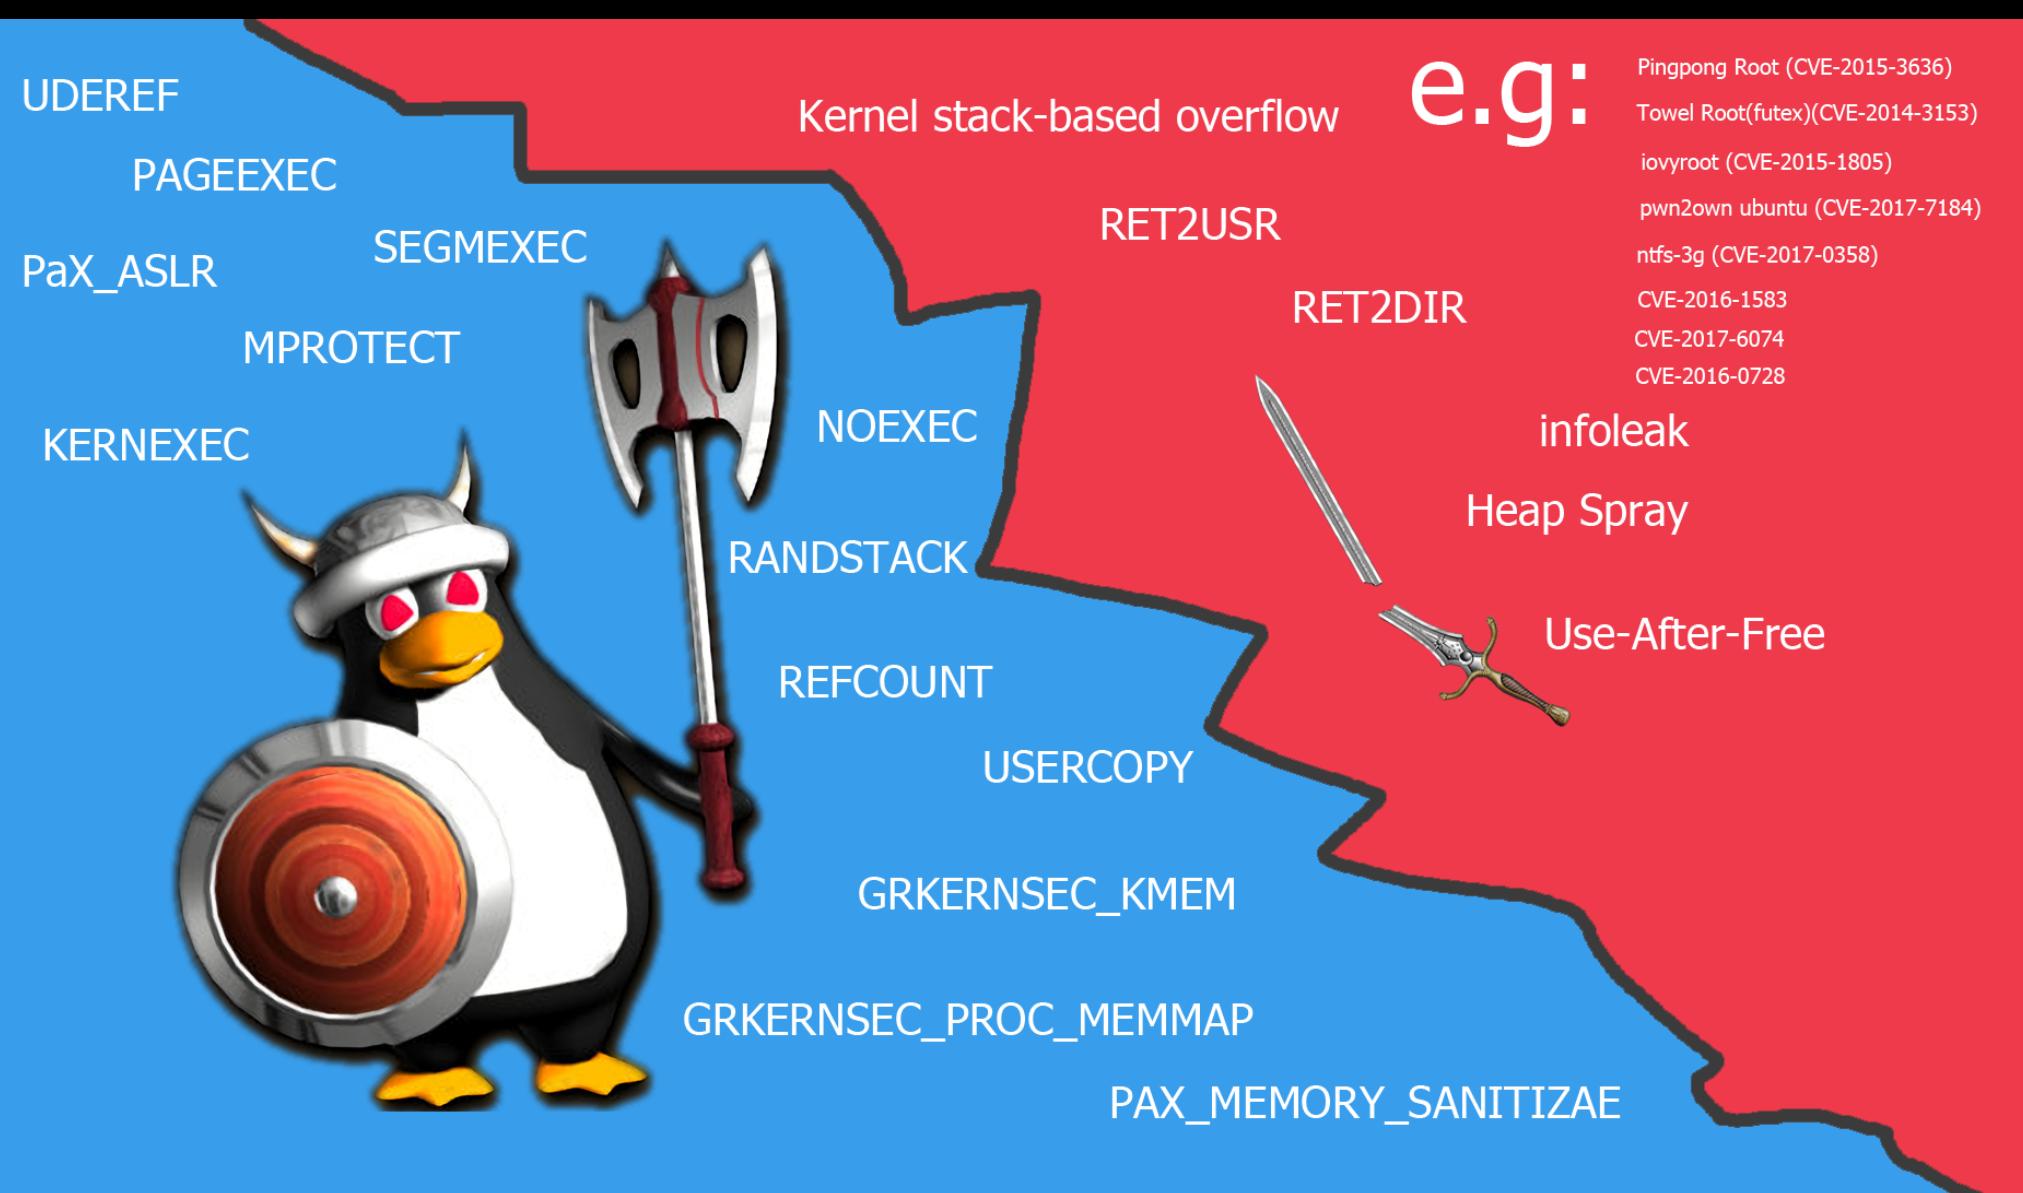
# #End of story?

We got everything we need! Ok, this is it.  
Thanks for coming. Bye, cruel world!

# #Why attack the core?

- \* Linux kernel sucks in 2000s
  - \* One null-ptr deref can rule them all
- \* Still a cargo-cult security in 2017?
- \* Harder to exploit userspace programs

# #Defend the core



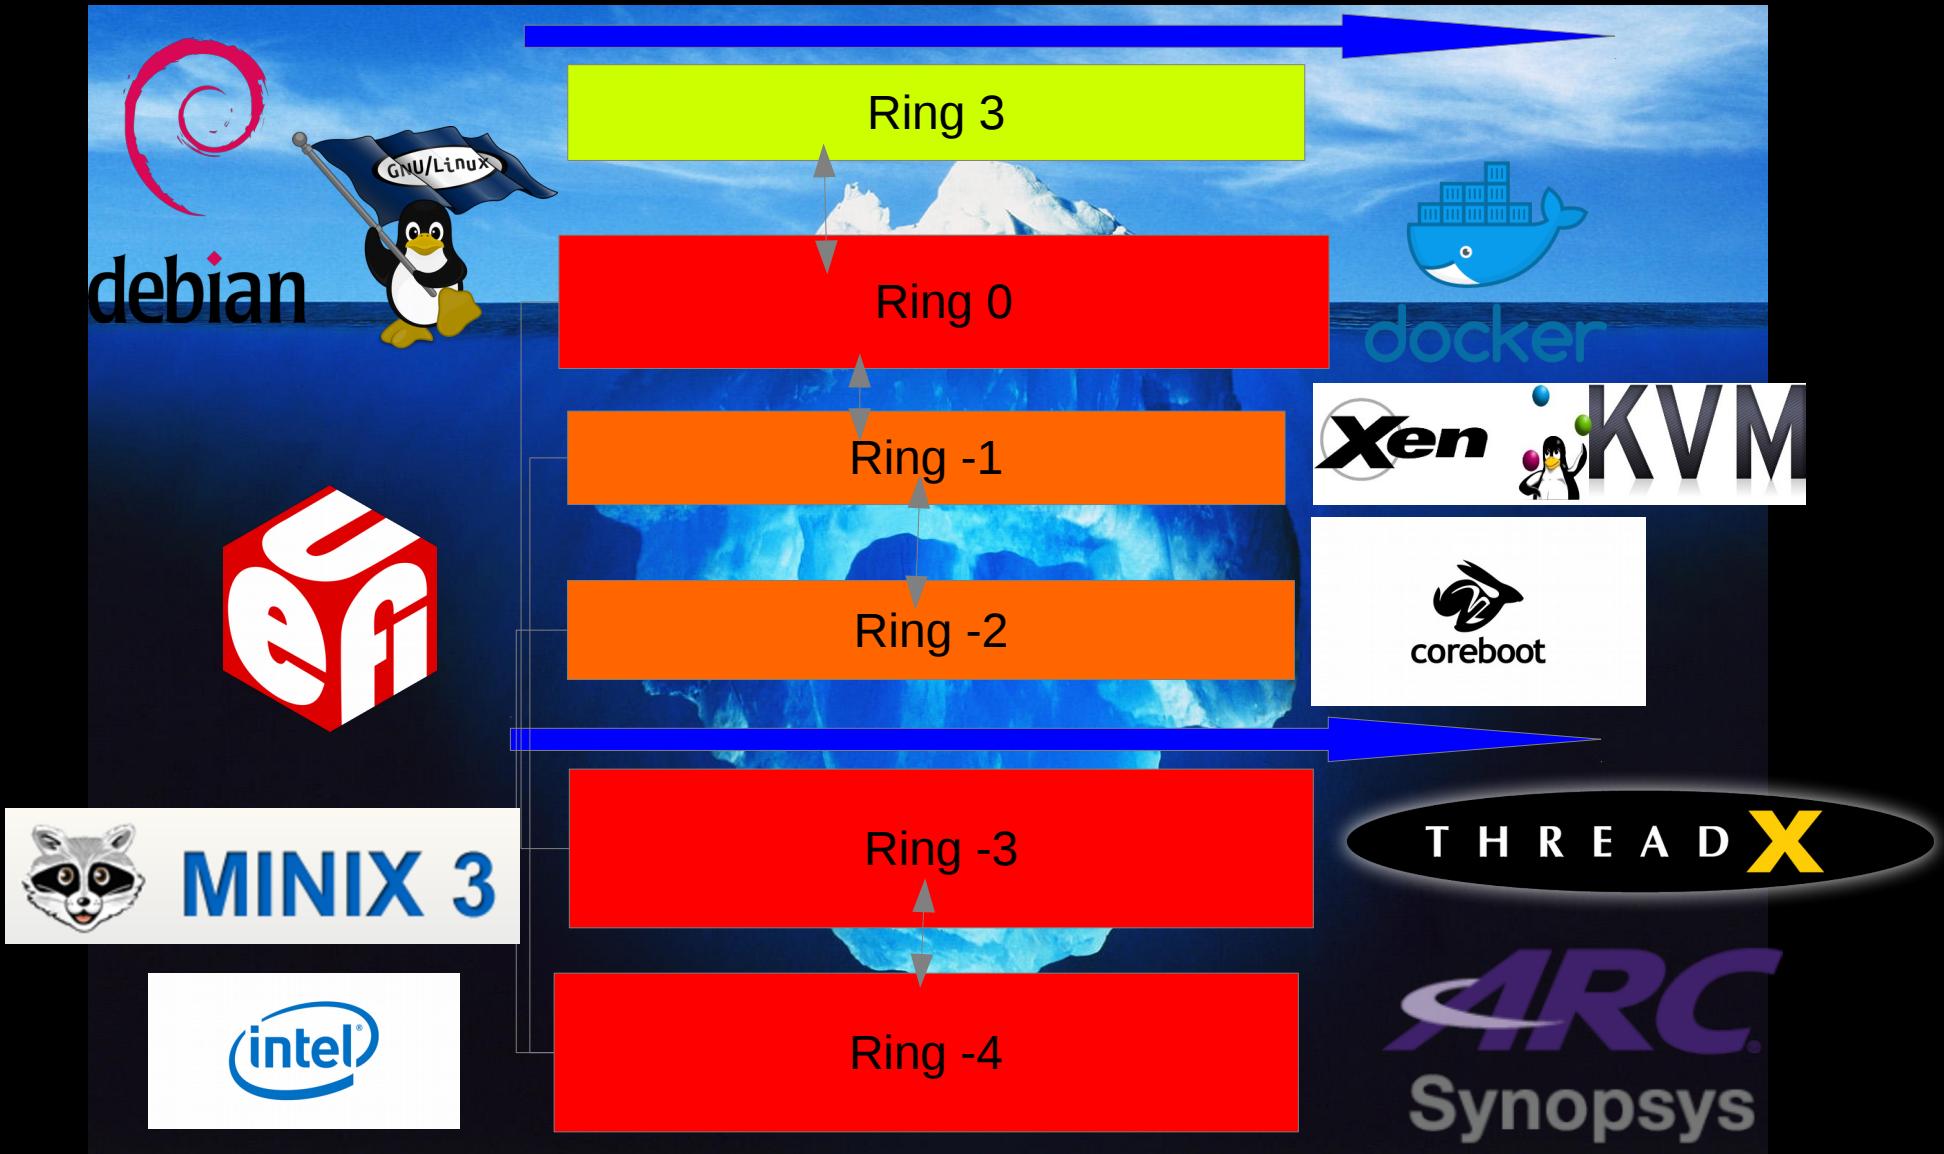
# #Statistics – 2016/2017

Sources	CVEs	Public exploit works?
MITRE	245	1
Ubuntu security tracker	105	1

Only one public exploit work at PaX/Grsecurity  
in two yrs, and you can kill that one w/  
situational hardening.

# #Wait, which CORE exactly?

\* RING 0 is **not** the CORE anymore



# #Under the water

- \* VM guest escape
- \* Did all device drivers follow the best practice( let's say IOMMU)?
- \* Old/new good/bad attacks on SMM
- \* Persistent attack chain
  - \* -1: Hijacking the kernel( perf impact?)
  - \* -2: Memory tricks

See the reference...

# #Demo

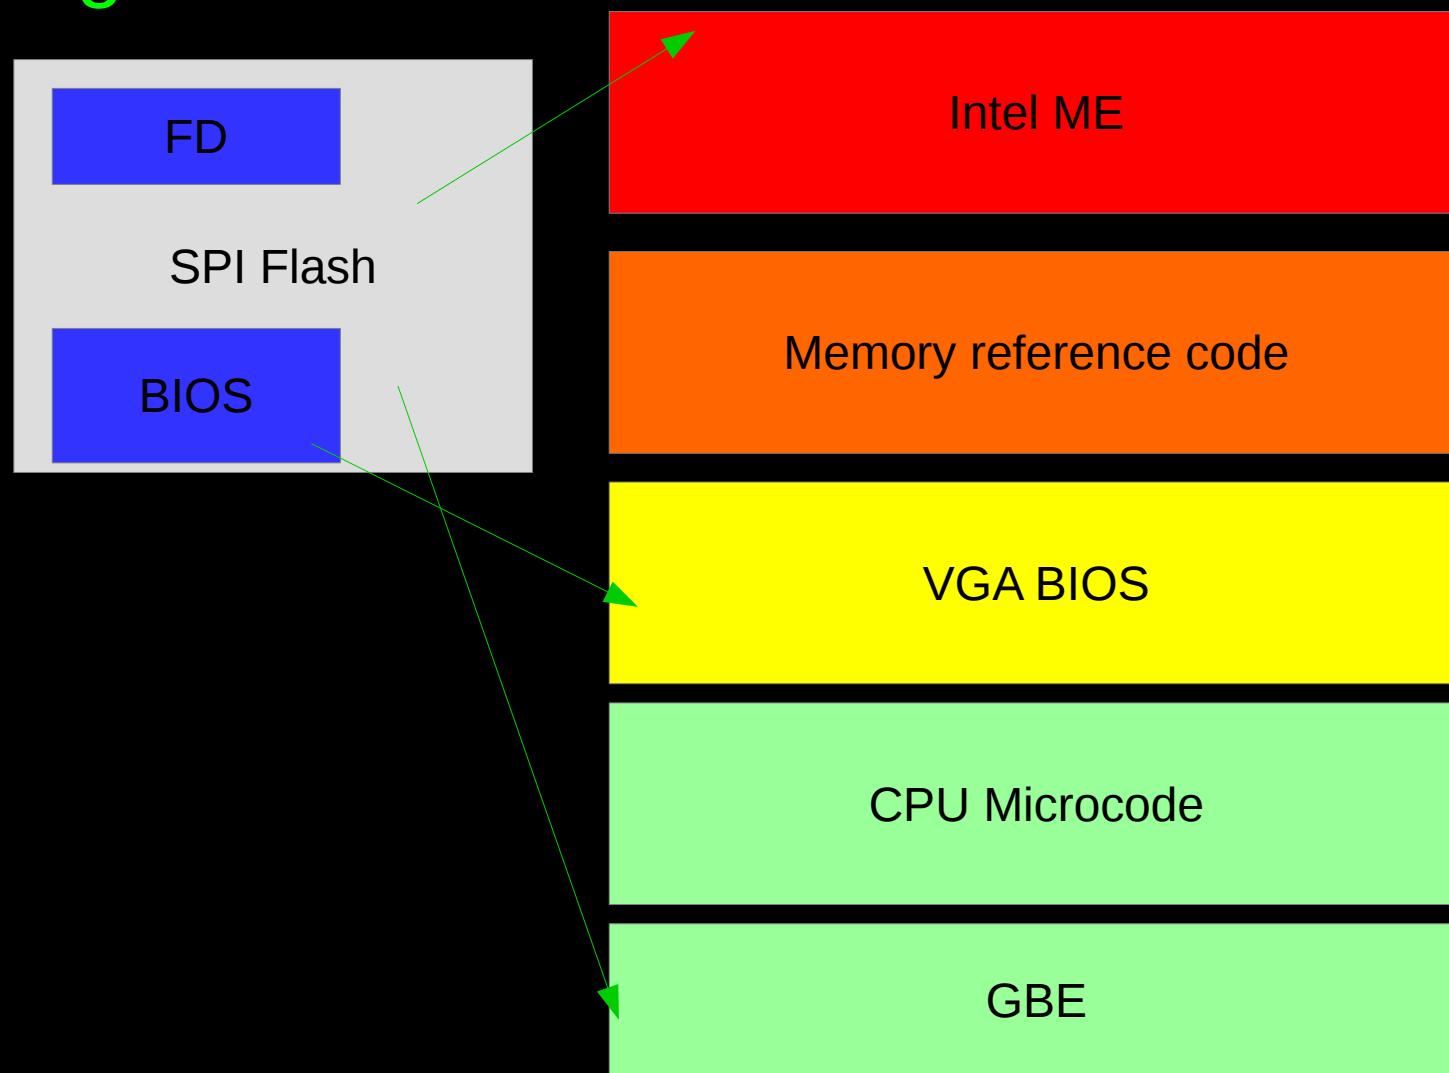
What could possibly go wrong with your  
“trusted” watcher? Don’t freaking out...

# #Invincible devil

- \* Intel ME
- \* A lot of ME "apps" based on it
- \* Changed a lot since v11
- \* Can't be disabled until HAP/altmedisable bit disclosed that we know how BIGBRO does about Intel ME for their own defense

# #From a libre FW's perspective

\* Not good



# #Hope or delusion?

- \* Some trade-off must be made
- \* Reproducible builds for PaX/Grsecurity
- \* Hardenedboot: Measured boot & verified boot
- \* Restricted ME via minimizing its functions

# #Why PaX/Grsecurity matters

- \* Kill bug classes
- \* Kill exploit vectors
- \* Assumption: Let data center takes physical security into account
  - \* Kernel is still the path to the under water
  - \* Hardened guards kills the attack surfaces

# #New attack surfaces

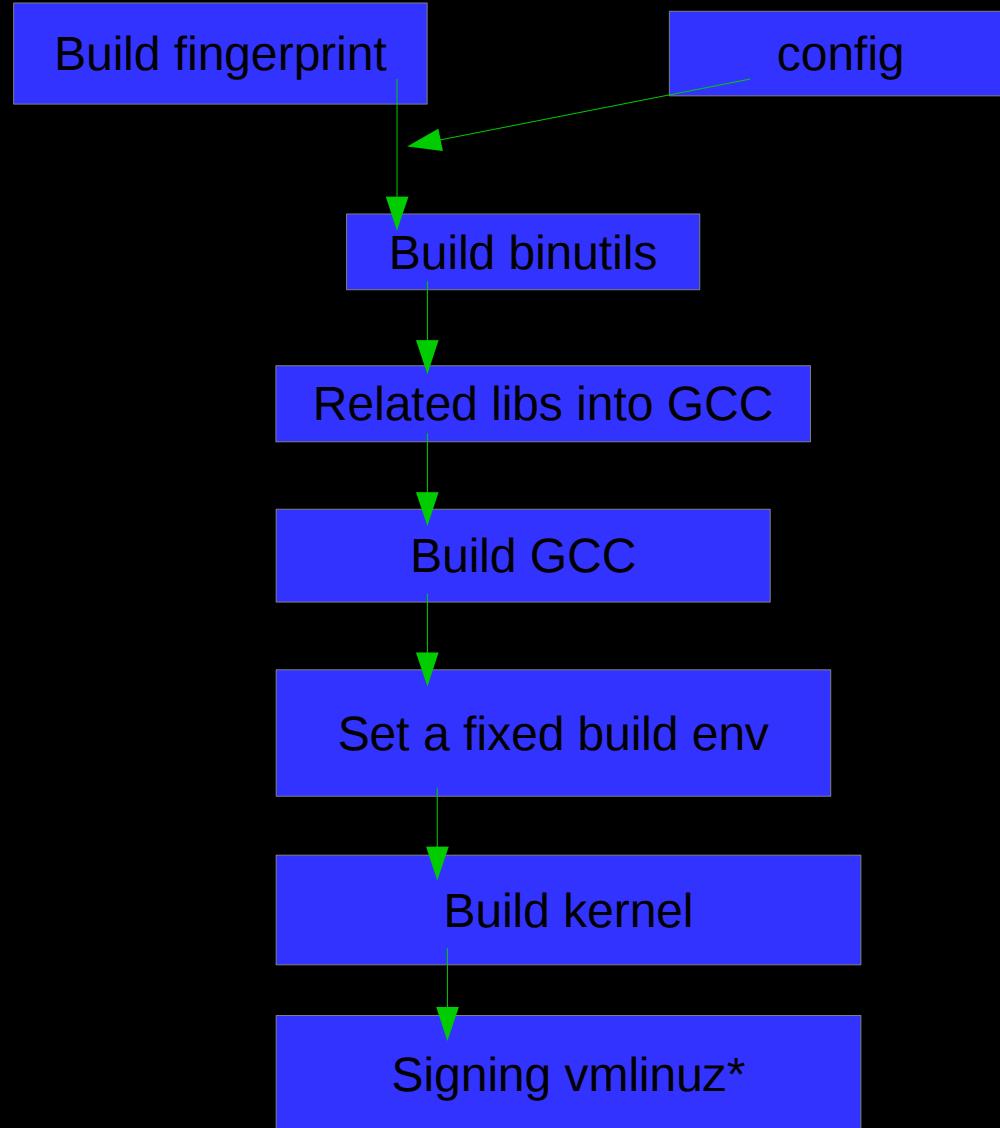
- \* Why Side-channel/Fault injection should be considered in “Hardening the COREs” solution?
- \* Sorry, “Raising the bar” isn’t what PaX/Grsecurity does, it never was.

```
01774 loc_1774:           ; CODE XREF: sub_1730+A8↓
01774     lea    rcx, [rbp+var_30]
01778     xor    r8d, r8d
0177B     mov    edx, ebx
0177D     mov    rsi, r12
01780     mov    edi, 0FFFFFFF9Ch
01785     jmp    short loc_1794
01785 ; -
01787     db    48h
01788     dq    OFFFFFFFFA614E6D8B8h
01790     db    0FFh, 3 dup(0CCh)
01794 ; -
01794 loc_1794:           ; CODE XREF: sub_1730+55↑
01794     call   $+
01799     test   eax, eax
0179B     jz    short loc_17DC
0179D loc_179D:           ; CODE XREF: sub_1730+E5↓
0179D     cmp    eax, 0xFFFFFFF8Ch
017A0     jz    short loc_17CC
017A2 loc_17A2:           ; CODE XREF: sub_1730+AA↓
017A2     cdqe
017A4 loc_17A4:           ; CODE XREF: sub_1730+EE↓
017A4     mov    rdx, [rbp+8]
017A8     cmp    qword ptr [rdx-10h], OFFFFFFFFFA51F98E0h
017B0     jnz    short loc_1825
017B2     mov    rsi, [rbp+var_20]
017B6     xor    rsi, gs:28h
017BF     jnz    short loc_1820
017C1     add    rsp, 20h
017C5     pop    rbx
017C6     pop    r12
017C8     pop    r13
017CA     pop    rbp
017CB     retn
```

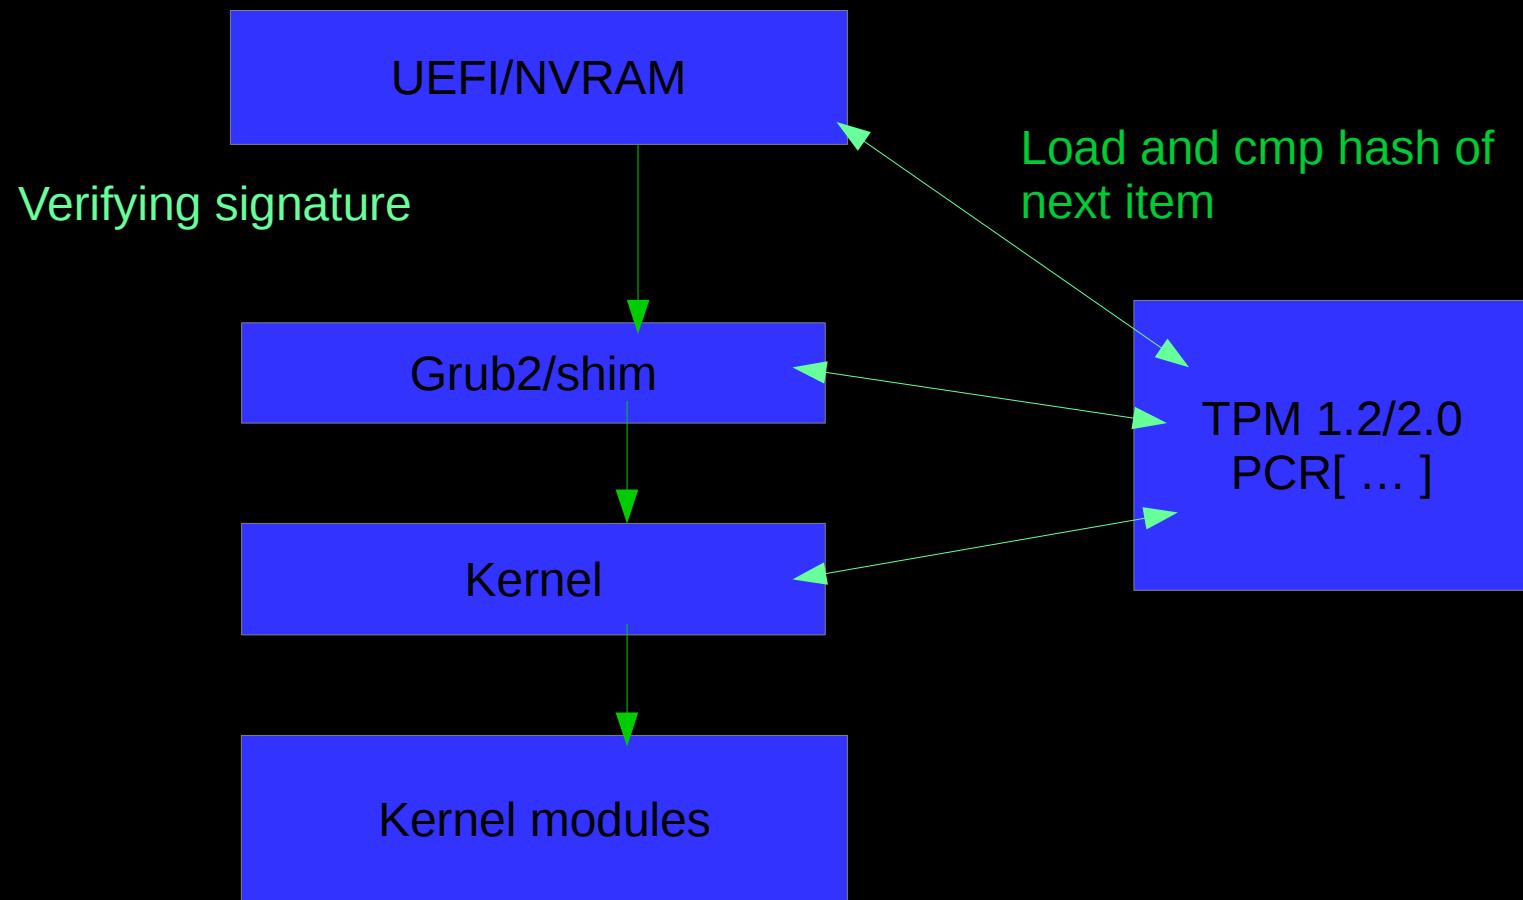
```
; arg int arg_8h @ rbp+0x8
; JMP XREF from 0x080051a6 (loc.08005195)
; CALL XREF from 0x080051b5 (loc.080051b5) ; [1]
0x080051b5 e890000000 callq 0x80051ba
; CALL XREF from 0x080051b5 (loc.080051b5) ; [0]
0x080051ba 48b5508 movq arg_8h,%rdx
0x080051be 48817af0c4e7. cmpq $-0x21c8183c,-0x10(%rdx)
0x080051c6 7507 if (var) goto 0x80051cf ; [2]
0x080051c8 5b popq %rbx
0x080051c9 415c popq %r12
0x080051cb 415d popq %r13
0x080051cd 5d
0x080051ce c3 retq
0x080051cf 0fb9 ud2b
0x080051d1 100f adcb %cl,(%rdi)
0x080051d3 1f invalid
0x080051d4 4000662e addb %spl,0x2e(%rsi) ; '0
0x080051d8 0f1f84000000. nopl (%rax,%rax)
0x080051e0 cc int3
```

```
{ v3 = 1;
do
{
    if ( !_DWORD(result) )
        LODWORD(result) = 0;
    if ( (_DWORD)result != -116 )
        break;
    v4 = v3 & 0x20;
    v3 = 33;
}
while ( !v4 );
result = (signed int)result;
}
if ( *(retaddr - 2) != -1524655904LL || __readgsqword(0x28u) != v5 )
    BUG();
return result;
```

# #Reproducible builds for PaX/Grsecurity



# #Hardened Boot



# #Free/libre solution?

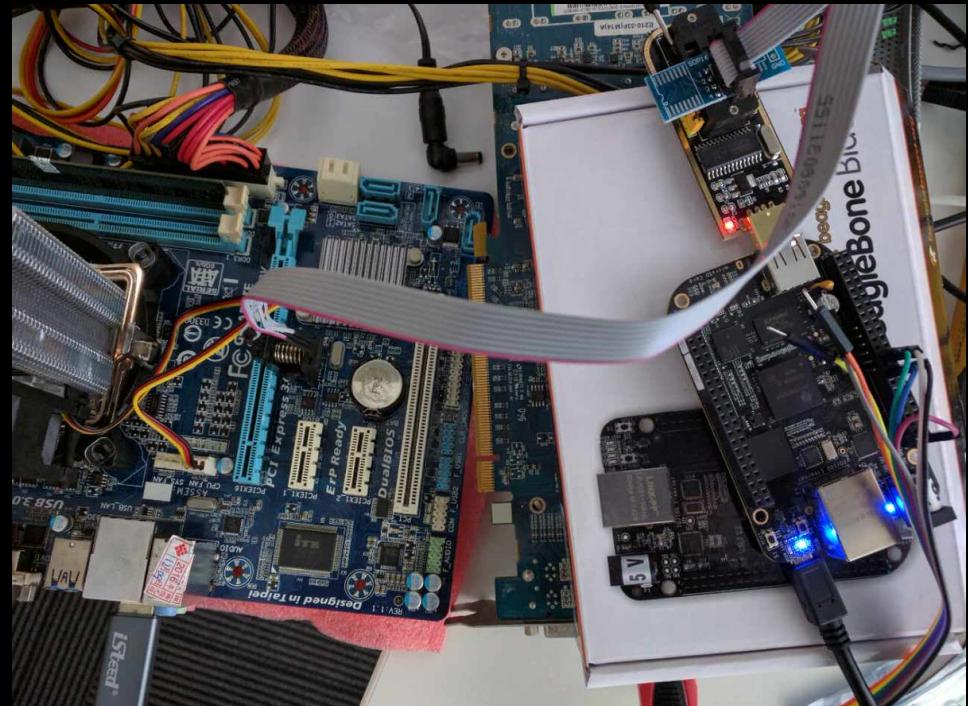
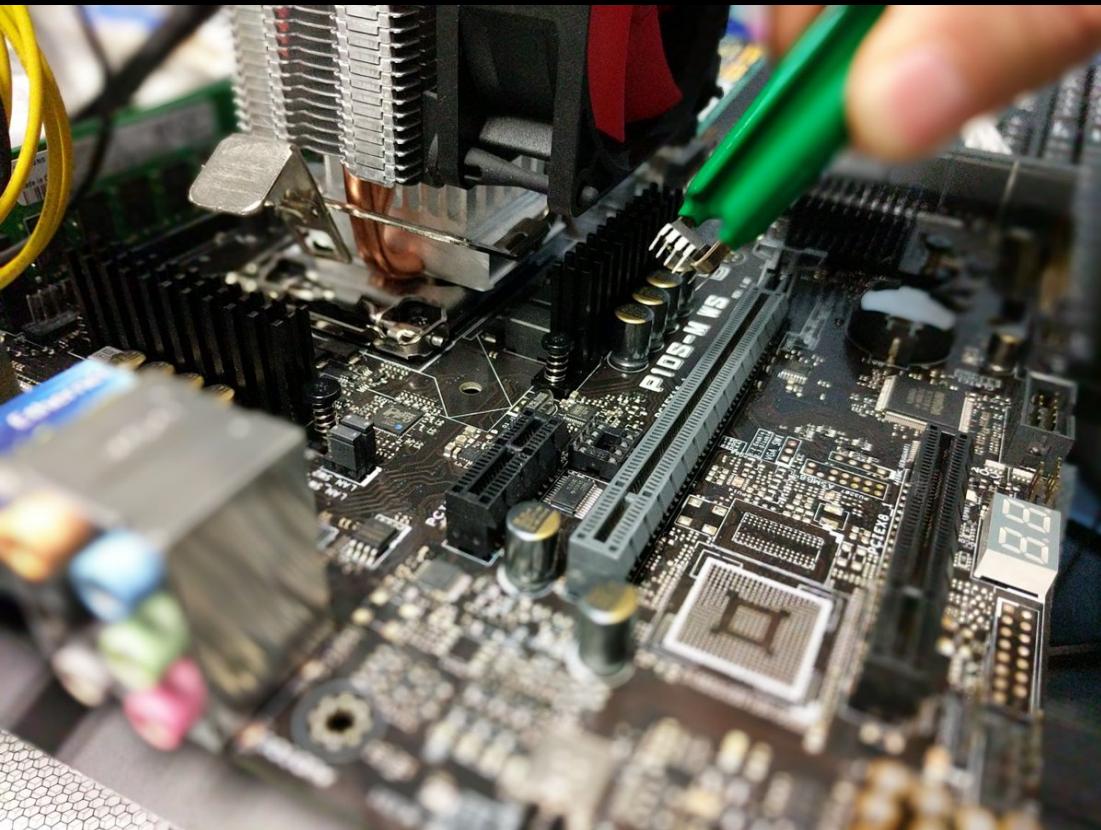
- \* Coreboot/FSP
- \* Libreboot
- \* Situational hardening
  - \* Old good machines for critical assets
  - \* Newer machines for generic purposes
  - \* Security needs some trade-off if performance is necessary: NERF/LinuxBoot
- \* Reference model for custom firmware
- \* Long-term plan: coreboot for RISC-V?

# #Open == auditable

\* open != secure, a-random-coreboot-machine:

```
[CHIPSEC] **** SUMMARY ****
[CHIPSEC] Time elapsed          0.014
[CHIPSEC] Modules total        17
[CHIPSEC] Modules failed to run 0:
[CHIPSEC] Modules passed       4:
[+] PASSED: chipsec.modules.common.smm
[+] PASSED: chipsec.modules.common.ia32cfg
[+] PASSED: chipsec.modules.common.smrr
[+] PASSED: chipsec.modules.common.spi_fdopss
[CHIPSEC] Modules failed       8:
[-] FAILED: chipsec.modules.common.bios_wp
[-] FAILED: chipsec.modules.common.bios_ts
[-] FAILED: chipsec.modules.common.spi_lock
[-] FAILED: chipsec.modules.common.spi_desc
[-] FAILED: chipsec.modules.common.bios_smi
[-] FAILED: chipsec.modules.memconfig
[-] FAILED: chipsec.modules.smm_dma
[-] FAILED: chipsec.modules.remap
[CHIPSEC] Modules with warnings 2:
[!] WARNING: chipsec.modules.common.bios_kbrd_buffer
[!] WARNING: chipsec.modules.common.rtclock
[CHIPSEC] Modules skipped 3:
[*] SKIPPED: chipsec.modules.common.uefi.s3bootscript
[*] SKIPPED: chipsec.modules.common.uefi.access_uefispec
[*] SKIPPED: chipsec.modules.common.secureboot.variables
[CHTPSFC] ****
```

# #Neutralizing ME



Mainboard	CPU	Tested BIOS
GA-B75M-D3H	SandyBridge	OEM/Coreboot
GA-B75M-D3V	IvyBridge	OEM/Coreboot
Lenovo T420	IvyBridge	OEM/Coreboot
Lenovo X220/X220i	SandyBridge	OEM/Coreboot
Lenovo X230	IvyBridge	OEM/Coreboot
Chromebook XE550C22	IvyBridge	OEM/Coreboot
ASUS P10S-M WS	Skylake	OEM

## #Extra notes;-)

- \* Neutralized ME affected remote attestation via SGX?
- \* Think harder about your situational hardening solution before ask OEM write your pk
- \* You don't need Intel Bootguard if the data center could take care of physical security
- \* Big clouds( AWS/Google/Facebook/BAT3H/etc) should do the situational hardening for their core infrastructure. Otherwise it might become someone else's computer;-)
- \* Masters of Pwn killed by PaX/Grsecurity

## #Are we done?

- \* Or just another starting point?
- \* Know your enemy( from Ring 3/0/-1/-2/-3/-4)
- \* Risk assessment for important production/assets
  - \* Known bug classes/exploit vectors

**HardenedLinux's  
roadmap**



The West Of  
Middle Earth  
At The End Of  
The Third Age

Crypto  
Engineering

Firmware

Compiler

Security  
operations

Situational  
Hardening

FORODWAITH  
Angmar  
Mount Gunduband  
Ered Mithrin  
Iron Hills  
HITHAEGLIR  
Ettenmoors  
Weather Hills  
Rhudaur  
Old Forest  
The Shire  
Bree  
Rivendell  
Eregion  
Dunland  
Isengard  
Lorien  
Fangorn  
Helms Deep  
N. Ithilien  
S. Ithilien  
Plateau of Gorgoroth  
Anfalas  
Lebennin  
RHŪN  
Mirkwood  
East Bight  
Erath on the Long Lake  
Sea of Rhun

Free/libre & Open  
Source Software's eco-  
system!

\* GPL-compliance  
\* Legislation  
\* Education

Adversary

# Reference

\* PaX/Grsecurity:

<https://grsecurity.net/>

\* Coreboot:

<http://coreboot.org/>

\* Linux kernel mitigation checklist:

[https://hardenedlinux.github.io/system-security/2016/12/13/kernel\\_mitigation\\_checklist.html](https://hardenedlinux.github.io/system-security/2016/12/13/kernel_mitigation_checklist.html)

\* Ring 0: Linux kernel vulnerability & exploitation & silent fixes

[https://github.com/hardenedlinux/grsecurity-101-tutorials/blob/master/kernel\\_vuln\\_exp.md](https://github.com/hardenedlinux/grsecurity-101-tutorials/blob/master/kernel_vuln_exp.md)

\* Virtualization security:

[https://github.com/hardenedlinux/grsecurity-101-tutorials/blob/master/virt\\_security.md](https://github.com/hardenedlinux/grsecurity-101-tutorials/blob/master/virt_security.md)

# Reference

\* Firmware security:

[https://github.com/hardenedlinux/firmware-anatomy/blob/master/hack\\_ME/firmware\\_security.md](https://github.com/hardenedlinux/firmware-anatomy/blob/master/hack_ME/firmware_security.md)

\* Reproducible builds for PaX/Grsecurity

<https://github.com/hardenedlinux/grsecurity-reproducible-build>

\* Hardened Boot:

[https://github.com/hardenedlinux/Debian-GNU-Linux-Profiles/tree/master/docs/hardened\\_boot](https://github.com/hardenedlinux/Debian-GNU-Linux-Profiles/tree/master/docs/hardened_boot)

\* Neutralized ME with coreboot stuff

[https://github.com/hardenedlinux/hardenedlinux\\_profiles/tree/master/coreboot](https://github.com/hardenedlinux/hardenedlinux_profiles/tree/master/coreboot)

\* Intel ME's info:

[https://github.com/hardenedlinux/firmware-anatomy/blob/master/hack\\_ME/me\\_info.md](https://github.com/hardenedlinux/firmware-anatomy/blob/master/hack_ME/me_info.md)

QA

Thanks